Prompt Library v1

AI-Powered LinkedIn Content Generation Assistant

GENERATIVE AI AND PROMPT ENGINEERING INTERNSHIP

TASK 2

Table of Contents

1. Introduction

Prompt Engineering is the discipline of designing, structuring, and refining the instructions given to a large language model (LLM) so that it reliably produces the intended output. Rather than treating a prompt as a single casual instruction, prompt engineering approaches it as a piece of reusable, testable design — one that specifies the model's role, the context of the task, what it should produce, how the output should be formatted, and what constraints it must respect.

Reusable prompt templates are important because they turn a one-off instruction into a repeatable asset. Instead of rewriting a prompt from scratch every time a similar task arises, a well-designed template can be parameterized with variables and reused across many scenarios, saving time and reducing inconsistency between outputs.

Structured prompts — those that clearly separate role, context, task, format, and constraints — improve the consistency of AI output because they remove ambiguity about what is expected. When the model is told exactly who it is acting as, what information it has, what it needs to produce, how the result should be organized, and what boundaries it must stay within, the variance between generated outputs drops significantly, even as the underlying topic changes.

In real-world Generative AI workflows, prompt libraries play the same role that reusable code libraries play in software engineering: they capture proven, tested instructions in a central place so that teams and individuals can apply them consistently across projects, onboard new use cases faster, and maintain a predictable quality bar across all AI-generated content.

2. Objective

The objective of this project is to design and demonstrate a reusable prompt library built around the Role + Context + Task + Format + Constraints framework. Under this framework, a single template defines a fixed structure — who the AI should act as, what background it should assume, what deliverable it should produce, how that deliverable should be formatted, and what rules it must follow — while a small set of variables (topic, audience, key points, tone, and length) are changed between uses.

This report demonstrates that principle in practice: one reusable template is used to generate five distinct, professional LinkedIn posts — spanning a project announcement, an automation workflow showcase, a product launch, an educational explainer, and a career journey post — simply by substituting the variables for each scenario while keeping the underlying template structure unchanged.

3. Reusable Prompt Template

The template below is the single reusable asset from which every prompt demonstration in Section 5 is derived. Only the variables inside the Task section change between uses; the Role, Context, Format, and Constraints remain fixed.

ROLE:

You are an expert LinkedIn content strategist specializing in AI, software engineering, and technology branding.

CONTEXT:

I need to create a professional LinkedIn post for a technology audience. The post should highlight technical achievements, projects, learning experiences, or industry insights while maintaining an engaging and professional tone.

TASK:

Create a LinkedIn post based on the provided details.

Variables:

- Topic: {{topic}}
- Target Audience: {{audience}}
- Key Points: {{key_points}}
- Tone: {{tone}}
- Length: {{length}}

FORMAT:

Generate the output with:

1. Attention-grabbing opening hook
2. Main content with clear paragraphs
3. Key technical points (if applicable)
4. Closing statement
5. Relevant hashtags

CONSTRAINTS:

- Avoid generic statements.
- Keep it professional and human-like.
- Do not exaggerate achievements.
- Use simple but impactful language.
- Maximum {{length}} words.

4. Prompt Demonstrations and AI Outputs

The five sections below apply the reusable template to five different content scenarios. Each section presents the complete prompt used, the corresponding AI-generated LinkedIn post, and a short explanation of the prompt's purpose.

Prompt 1: Generative AI Project Announcement

Complete Prompt

Role:

You are an expert LinkedIn content strategist specializing in AI, software engineering, and technology branding.

Context:

I need to create a professional LinkedIn post for a technology audience. The post should highlight technical achievements, projects, learning experiences, or industry insights while maintaining an engaging and professional tone.

Task:

Create a LinkedIn post based on the following details:

Topic:

Building my first Generative AI application

Target Audience:

AI Engineers, recruiters, and technology professionals

Key Points:

LLM API integration, Python development, prompt engineering, real-world AI application, learning experience

Tone:

Professional and reflective

Length:

200 words

Format:

Generate:

1. Attention-grabbing opening hook
2. Main content with clear paragraphs
3. Key technical points
4. Closing statement
5. Relevant hashtags

Constraints:

- Avoid generic statements.
- Keep it professional and human-like.
- Do not exaggerate achievements.
- Use simple but impactful language.

- Maximum 200 words.

Prompt 1 Output — Generative AI Project Announcement

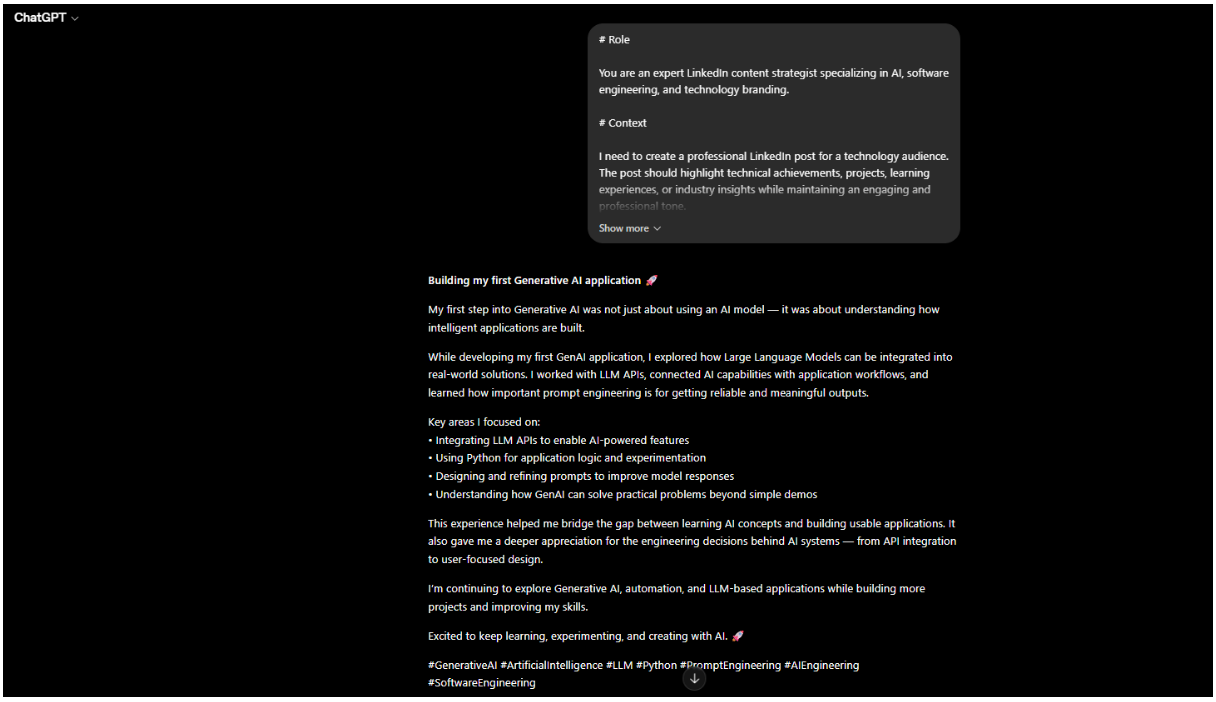

Figure 1 — AI-generated LinkedIn post for the Generative AI Project Announcement prompt.

Purpose

This prompt was designed to help a technology professional announce a hands-on Generative AI project in a way that reads as authentic and reflective rather than promotional. It anchors the model to concrete technical details — LLM API integration, Python, and prompt engineering — so the generated post highlights real skills instead of vague enthusiasm.

Prompt 2: AI Automation Workflow Showcase

Complete Prompt

Role:

You are an expert LinkedIn content strategist specializing in AI, automation, and emerging technologies.

Context:

I need to create a professional LinkedIn post showcasing a technical AI automation project for a technology audience.

Task:

Create a LinkedIn post based on the following details:

Topic:

Developing AI-powered automation workflows

Target Audience:

AI professionals, software developers, and business technology leaders

Key Points:

n8n workflow automation, API integrations, AI agents, reducing repetitive tasks, improving business processes

Tone:

Professional and educational

Length:

180 words

Format:

Generate:

1. Strong opening hook
2. Project overview
3. Technical implementation details
4. Business impact
5. Closing statement with hashtags

Constraints:

- Avoid generic AI buzzwords.
- Maintain a professional tone.
- Explain technical concepts clearly.

- Do not overstate results.
- Maximum 180 words.

Prompt 2 Output — AI Automation Workflow Showcase

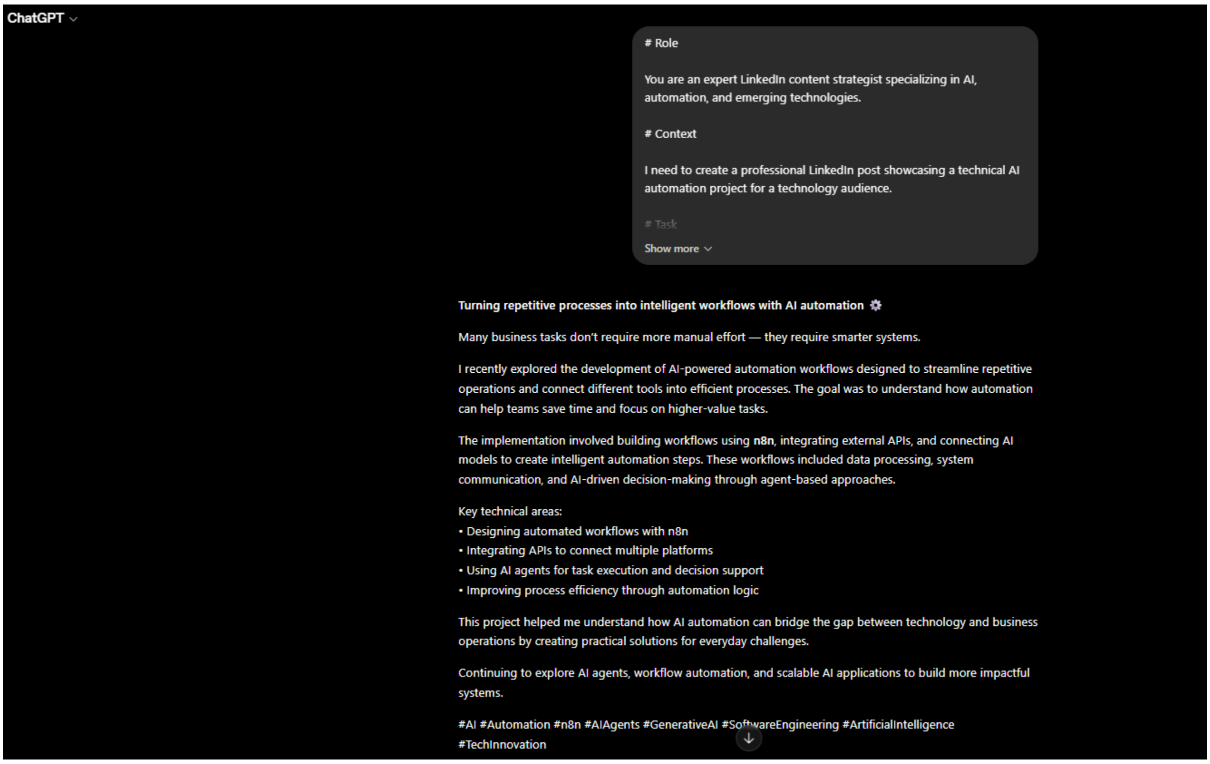

Figure 2 — AI-generated LinkedIn post for the AI Automation Workflow Showcase prompt.

Purpose

This prompt targets a more technically literate audience and asks the model to explain an automation workflow (built with n8n and API integrations) in terms of both implementation and business impact. It is useful for demonstrating how a single template can shift tone from reflective storytelling to educational, results-oriented writing simply by changing the variables.

Prompt 3: AI Product Launch Announcement

Complete Prompt

Role:

You are a LinkedIn content strategist for AI startups and technology companies.

Context:

I need to create an engaging LinkedIn post announcing an AI-powered product while highlighting its technical innovation and user impact.

Task:

Create a LinkedIn post based on the following details:

Topic:

Launching an AI-powered language learning platform

Target Audience:

Startup founders, developers, AI innovators, and technology professionals

Key Points:

Conversational AI, LLM integration, speech recognition, text-to-speech, personalized learning experience

Tone:

Achievement-focused and professional

Length:

200 words

Format:

Generate:

1. Engaging introduction
2. Product explanation
3. Technical highlights
4. User benefits
5. Professional closing and hashtags

Constraints:

- Keep the content authentic.
- Avoid exaggerated claims.

- Maintain a technology-focused tone.
- Maximum 200 words.

Prompt 3 Output — AI Product Launch Announcement

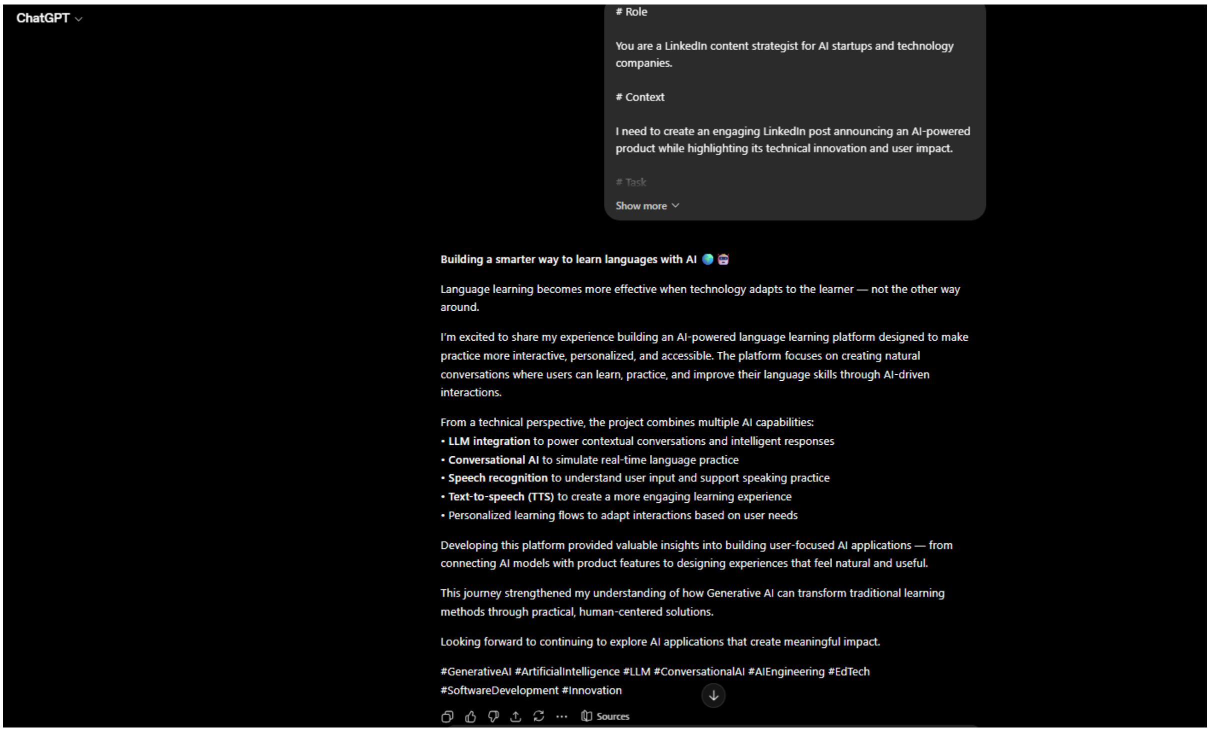

Figure 3 — AI-generated LinkedIn post for the AI Product Launch Announcement prompt.

Purpose

This prompt reframes the same template for a product-launch use case, balancing excitement about a new AI-powered platform with authenticity constraints that prevent overstated marketing claims. It shows the template's flexibility in producing startup-style announcement content without losing a grounded, professional tone.

Prompt 4: AI Educational Content Creation

Complete Prompt

Role:

You are an AI educator and technical content creator specializing in explaining complex AI concepts.

Context:

I need to create an educational LinkedIn post that explains an AI concept in a simple and professional way.

Task:

Create a LinkedIn post based on the following details:

Topic:

Understanding Retrieval-Augmented Generation (RAG) systems

Target Audience:

Computer science students, AI beginners, and technology professionals

Key Points:

RAG architecture, embeddings, vector databases, external knowledge retrieval, improving LLM responses

Tone:

Educational and beginner-friendly

Length:

250 words

Format:

Generate:

1. Interesting opening question or hook
2. Simple explanation of the concept
3. Technical breakdown
4. Practical applications
5. Learning takeaway and hashtags

Constraints:

- Explain concepts clearly.
- Avoid unnecessary complexity.
- Maintain technical accuracy.

- Maximum 250 words.

Prompt 4 Output — AI Educational Content Creation

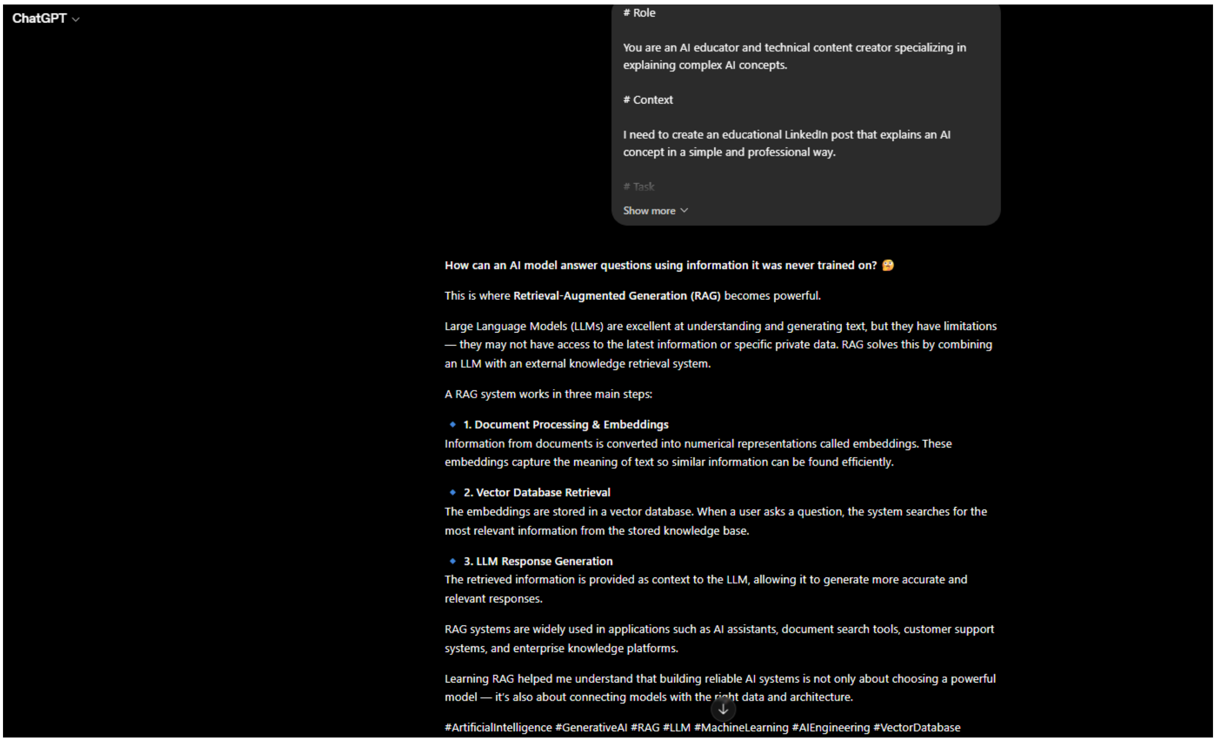

Figure 4 — AI-generated LinkedIn post for the AI Educational Content Creation prompt.

Purpose

This prompt repurposes the template for educational rather than self-promotional content, asking the model to explain Retrieval-Augmented Generation to a beginner audience. It demonstrates how the Role and Tone variables alone can shift the same underlying structure from a personal-branding post into a technical explainer.

Prompt 5: AI Engineer Career Journey Post

Complete Prompt

Role:

You are a LinkedIn career content strategist specializing in technology and artificial intelligence.

Context:

I need to create a professional LinkedIn post about an AI learning journey that inspires students and aspiring engineers.

Task:

Create a LinkedIn post based on the following details:

Topic:

Roadmap towards becoming an AI Engineer

Target Audience:

Computer science students, aspiring AI engineers, and technology learners

Key Points:

Python programming, Generative AI, LLM applications, AI automation, building practical projects, continuous learning

Tone:

Motivational but technical

Length:

200 words

Format:

Generate:

1. Personal opening statement
2. Learning journey highlights
3. Technical skills discussion
4. Future goals
5. Closing reflection with hashtags

Constraints:

- Keep it realistic and authentic.
- Avoid motivational clichés.
- Maintain a professional technology-focused style.

- Maximum 200 words.

Prompt 5 Output — AI Engineer Career Journey Post

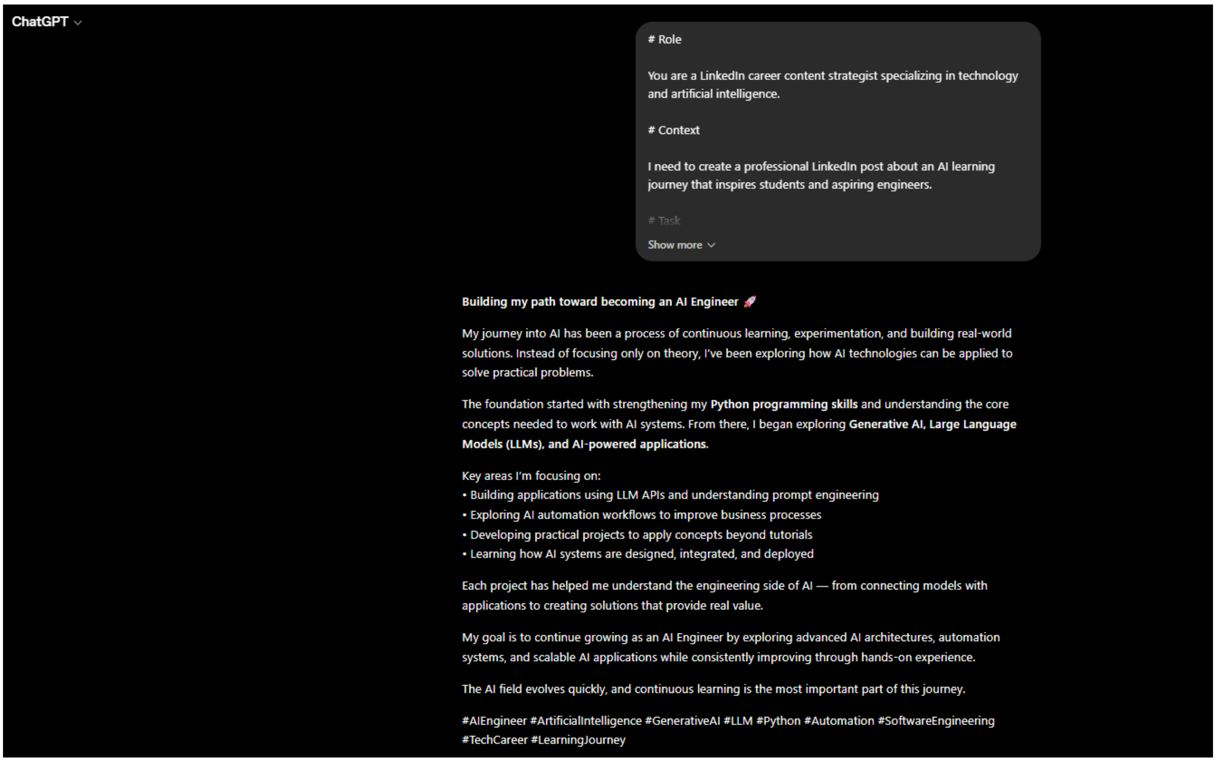

Figure 5 — AI-generated LinkedIn post for the AI Engineer Career Journey Post prompt.

Purpose

This final prompt applies the template to a first-person career narrative, balancing motivational framing with concrete technical grounding (Python, Generative AI, LLM applications). It rounds out the library by showing the template can support introspective, career-oriented content in addition to project and product announcements.

5. Key Learnings

- Reusable prompts reduce repeated prompt creation, since a single validated structure can be applied to many tasks instead of being rewritten each time.

- Variables make prompts adaptable for different scenarios, allowing the same template to generate content for announcements, showcases, launches, education, and career storytelling.
- Clear constraints improve AI output quality by preventing generic, exaggerated, or overly long responses and keeping the tone professional and human-like.
- Structured prompting — separating Role, Context, Task, Format, and Constraints — improves consistency, producing outputs with a predictable shape regardless of topic.
- Prompt libraries are useful for scalable AI workflows, since they let individuals and teams maintain a growing, reusable set of tested prompts rather than starting from scratch for every new task.

6. Conclusion

This project demonstrates that a single, well-structured prompt template — built on the Role + Context + Task + Format + Constraints framework — can reliably generate five distinct, professional pieces of content simply by changing a small set of variables. This reusability is central to practical Generative AI work: it reduces the effort required to produce consistent, high-quality output, makes prompt design auditable and improvable over time, and lays the foundation for building larger, scalable prompt libraries that can support real-world content, engineering, and product workflows.